

# Homework 7

Due May 28th, 2021 by 11:59pm

**Instructions:** This homework consists of a reading assignment and two coding exercises. Please submit your solutions via Gradescope. Solutions should consist of three files: a PDF containing your solutions to the *non-coding questions*; a Jupyter notebook (.ipynb file) and an html print-out (.html file) with your solution to the *coding exercise*. **All coding exercises must be completed in Python.** Please be sure to comment the code appropriately. Students are encouraged to discuss homework problems, particularly on Canvas and in the TA hours, but must submit their own solutions.

## Reading Assignment

- Review Lecture 7.
- Review Lab 7 (available from the Course Materials page on Canvas).

## Exercise 1

In this exercise you will implement your own version of a **kernel support vector machine** with the squared hinge loss. The kernel support vector machine with the squared hinge loss writes as

$$\min_{\alpha \in \mathbb{R}^n} F(\alpha) := \frac{1}{n} \sum_{i=1}^n \ell(y_i, (K\alpha)_i) + \lambda \alpha^T K \alpha, \quad (1)$$

where  $(K\alpha)_i$  is the  $i$ th entry in the vector  $K\alpha$ , and

$$\ell(y, t) := (\max\{0, 1 - yt\})^2. \quad (2)$$

**Note:** Despite the piecewise nature of this loss function, we claim that the quantity

$$\frac{1}{n} \sum_{i=1}^n \ell(y_i, (K\alpha)_i)$$

is differentiable at every  $\alpha \in \mathbb{R}^n$ . For reference on computing the gradient of a piecewise function (and confirming that this gradient exists everywhere), please refer to the **math review notes** on the course website.

- (a) Compute the gradient  $\nabla F(\alpha)$  of  $F$ .

- (b) The polynomial kernel of order  $p$  is given by

$$k(x, y) = (x^T y + b)^p.$$

Its parameters are the offset  $b$  and order  $p$ . Implement this kernel function.

- (c) Write a function *computegram* that computes, for any kernel  $k$  and set of datapoints  $x_1, \dots, x_n$ , the kernel matrix  $K$  with  $(i, j)^{th}$  entry  $k(x_i, x_j)$ .
- (d) Write a function *kerneleval* that computes, for any kernel  $k$ , set of datapoints  $x_1, \dots, x_n$  and a new datapoint  $x^*$ , the vector of kernel evaluations  $[k(x_1, x^*), \dots, k(x_n, x^*)]^T$ .
- (e) Write a function *mysvm* that implements the fast gradient algorithm to train the kernel support vector machine with the squared hinge loss. The function takes as input the initial step-size value for the backtracking rule and a stopping criterion based on the norm of the gradient.
- (f) Consider the Digits dataset ([http://scikit-learn.org/stable/modules/generated/sklearn.datasets.load\\_digits.html](http://scikit-learn.org/stable/modules/generated/sklearn.datasets.load_digits.html)). Download the data. Re-shape the images as vectors. Normalize each image vector such that it has norm 1; note that this is *not* the same as standardizing the data. You may use the function `normalize` from `sklearn.preprocessing`. Perform an 80-20 train-test split of the data.
- (g) For each digit  $d = 0, \dots, 9$  in the Digits dataset, train a kernel SVM classifier using the data  $X$  and binary label  $y^{(d)}$ , defined as

$$y_i^{(d)} = \begin{cases} 1 & y_i = d \\ -1 & y_i \neq d \end{cases}.$$

In this way, we obtain 10 “one vs rest” classifiers, each of which predicts whether a given image comes from class  $d$  or not. Use the polynomial kernel with  $p = 7$  and  $b = 1$ . Set  $\lambda = 10.0$ . For each classifier trained, plot the objective function versus the iteration number.

- (h) Now use 5-fold cross-validation to select the best choice of  $\lambda$  for each classifier trained above. Report the value of  $\lambda$  selected for each classifier. Re-train each classifier with its corresponding best choice of  $\lambda$ .
- (i) Recall that the SVM prediction is given by

$$\hat{f}(x) = \sum_{i=1}^n \alpha_i k(x, x_i).$$

In the multi-class setting, define the class prediction to be given by classifier that maximizes this score, out of the  $d$  “one vs rest” classifiers that you have trained. Generate test set predictions and report your final misclassification error.

## Exercise 2

Expand on the above analysis in an Azure virtual machine by comparing the performance of the following kernel functions:

- Linear:

$$k(x, y) = x^T y.$$

- Polynomial:

$$k(x, y) = (x^T y + b)^p$$

Try a few settings for the polynomial order  $p$ . Set either  $b = 1$  or  $b = 0$ .

- Gaussian / radial basis function (RBF):

$$k(x, y) = \exp(-\gamma \|x - y\|^2)$$

Try a few settings of the scale parameter  $\gamma$ .

For each kernel investigated, report your final misclassification error on the test set.

As in the previous submission, please include the following: (1) a screenshot of your running VM in the Azure dashboard; (2) a screenshot of your command line after you `ssh` to the VM; (3) a screenshot of a running Python notebook on your VM, where you have run the following code in a cell block:

```
import os
os.getcwd()
```

The screenshots may either be embedded in your notebook or submitted as additional files.